

Camera Models


Computer Vision – Lecture 14

(Generative Models postponed to Lecture 16)

Further Reading


- 7 lectures from [S Lazebnik](#)
- Slides from [D Fouhey and J Johnson](#)
- Many slides adapted from both sources

Multi-view geometry “Bible”



Perspective in Art

- Eastern art: planar perspective 11th century.
- Western art: early Renaissance: 15th century
- Before: attempts to simulate perspective but no good understanding.




Just look at this 14th-century painting of Jesus having his fateful tea party.

Things aren't always as they appear...



Single-view ambiguity




Single-view ambiguity



[Rashad Alakbarov shadow sculptures](#)

Anamorphic perspective




Anamorphic perspective




H. Holbein The Younger, *The Ambassadors*, 1533




Camera Obscura



Pinhole Camera




Perspective Projection



Camera Coordinate System


- The optical center (O) is at the origin
- The z axis is the *optical axis* perpendicular to the image plane
- The xy plane is parallel to the image plane, x and y axes are horizontal and vertical directions of the image plane

Perspective Projection




Perspective Projection

Instead of dealing with an image that is upside down, most of the time we will pretend that the image plane is *in front* of the camera center




Single View Ambiguity

From a single image, there is an inherent ambiguity about the scale of the observed object.




Projection of lines



Piero della Francesca, *Flagellation of Christ*, 1455-1460


Projection of lines

- Parallel lines meet at a *vanishing point*




Piero della Francesca, *Flagellation of Christ*, 1455-1460

Perspective Cues




Projection of 3D shapes

What is the shape
of the projection
of a sphere?



Projection of 3D shapes

- What is the shape of the projection of a sphere?



Homogeneous Coordinates III

To form homogeneous coordinates from normal Euclidean coordinates, append 1 as the last entry:

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous *image*
coordinates

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

homogeneous *scene*
coordinates

To convert *from* homogeneous coordinates, divide by the last entry:

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

$$\begin{bmatrix} x \\ y \\ z \\ w \end{bmatrix} \Rightarrow (x/w, y/w, z/w)$$

In homogeneous coordinates, all scalar multiples represent the same point!

Perspective Projection Matrix

- Projection is a matrix multiplication using homogeneous coordinates:


$$\begin{bmatrix} \\ \\ \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} \\ \\ \end{bmatrix} \Rightarrow \left(f \frac{x}{z}, f \frac{y}{z} \right)$$

divide by the third
coordinate

Orthographic projection

Special case of perspective projection


- Distance from center of projection to image plane is infinite
- Also called “parallel projection”



Orthographic projection

Special case of perspective projection


- Distance from center of projection to image plane is infinite
- Also called “parallel projection”



Orthographic projection

Special case of perspective projection

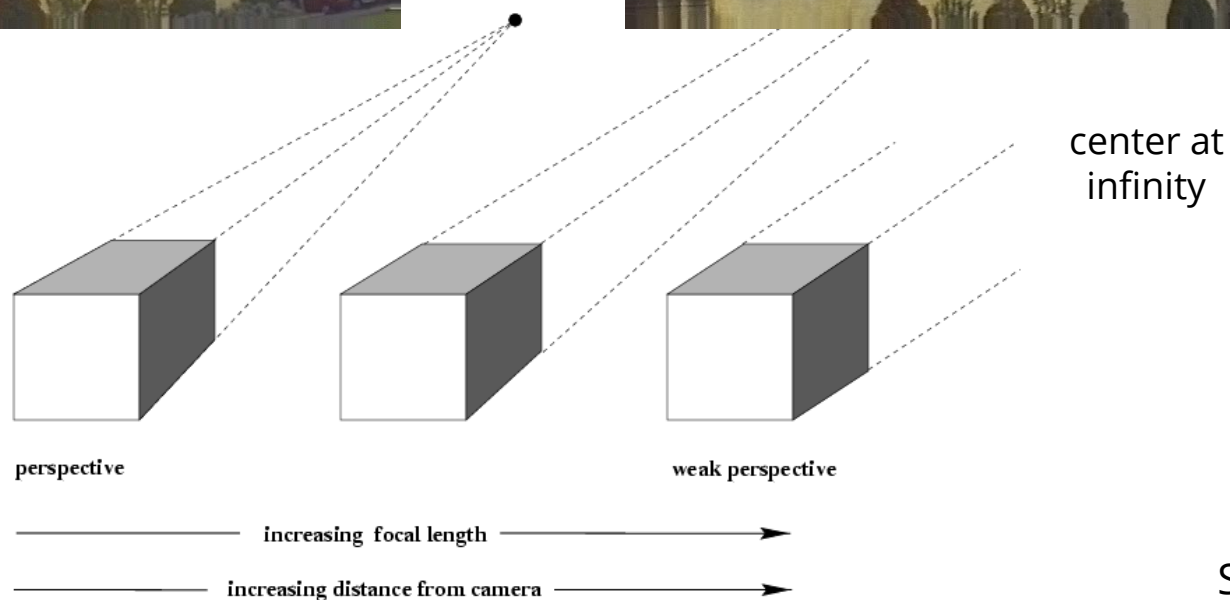
- Distance from center of projection to image plane is infinite
- Also called “parallel projection”




- Assuming projection along the z axis, what's the matrix?

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Approximating an orthographic camera





Normalized Coordinates



Normalized (camera) coordinate system: camera center is at the origin, the *principal axis* is the z -axis, x and y axes of the image plane are parallel to x and y axes of the world

Perspective projection in normalized coordinates



$$x = f \frac{X}{Z}, y = f \frac{Y}{Z}$$

$$\begin{pmatrix} x \\ y \\ 1 \end{pmatrix} \cong \begin{pmatrix} fX \\ fY \\ Z \end{pmatrix} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

Homogeneous coord.
vec. \mathbf{x} of image point


Camera
projection
matrix \mathbf{P}

Homogeneous coord.
vec. \mathbf{X} of 3D point


$$\mathbf{x} \cong \mathbf{P}\mathbf{X}$$

Equality up to scale

Camera Calibration



world coordinate system




Camera calibration: figuring out transformation from *world* coordinate system to *image* coordinate system

$$\begin{pmatrix} \text{2D point} \\ \mathbf{x} \\ (3 \times 1) \end{pmatrix} \cong \begin{pmatrix} \text{Camera to pixel coord. trans. matrix} \\ \mathbf{K} (3 \times 3) \end{pmatrix} \begin{pmatrix} \text{Canonical projection matrix} \\ [\mathbf{I} \mid \mathbf{0}] (3 \times 4) \end{pmatrix} \begin{pmatrix} \text{World to camera coord. trans. matrix} \\ \begin{bmatrix} \mathbf{R} & \mathbf{t} \\ \mathbf{0}^T & 1 \end{bmatrix} (4 \times 4) \end{pmatrix} \begin{pmatrix} \text{3D point} \\ \mathbf{X} \\ (4 \times 1) \end{pmatrix}$$


Intrinsic camera parameters:
principal point,
scaling factors

Extrinsic camera parameters:
rotation, translation

Camera Calibration



world coordinate system




Camera calibration: figuring out transformation from *world* coordinate system to *image* coordinate system

$$\begin{pmatrix} \text{2D point} \\ \mathbf{x} \\ (3 \times 1) \end{pmatrix} \cong \underbrace{\begin{pmatrix} \text{Camera to pixel coord. trans. matrix} \\ \mathbf{K} (3 \times 3) \end{pmatrix} \begin{pmatrix} \text{Canonical projection matrix} \\ [\mathbf{I} \mid \mathbf{0}] (3 \times 4) \end{pmatrix} \begin{pmatrix} \text{World to camera coord. trans. matrix} \\ \begin{bmatrix} \mathbf{R} & \mathbf{t} \\ \mathbf{0}^T & 1 \end{bmatrix} (4 \times 4) \end{pmatrix}}_{\mathbf{P} = \mathbf{K}[\mathbf{R}|\mathbf{t}]} \begin{pmatrix} \text{3D point} \\ \mathbf{X} \\ (4 \times 1) \end{pmatrix}$$

General camera projection matrix


Intrinsic parameters: Principal point



Principal point (p): point where principal axis intersects the image plane

- In the *normalized* coordinate system, the **origin** of the image is at the **principal point**
- In the *image* coordinate system: the **origin** is in the **corner**

Intrinsic parameters: Principal point




We want the principal point to map to (p_x, p_y) instead of $(0,0)$

$$x = f \frac{X}{Z} + p_x, \quad y = f \frac{Y}{Z} + p_y$$

$$\begin{pmatrix} x \\ y \\ 1 \end{pmatrix} \cong \begin{pmatrix} fX + Zp_x \\ fY + Zp_y \\ Z \end{pmatrix} = \begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

Intrinsic parameters: Principal point



Principal point: (p_x, p_y)

$$\begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} =$$

calibration
matrix K

Canonical
projection
matrix $[I \mid \mathbf{0}]$

$$\begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

$$P = K[I \mid \mathbf{0}]$$

Intrinsic parameters: Principal point

- What are the units of the focal length f and principal point coordinates (p_x, p_y) ?
 - Same as world units – likely metric units
- What units do we want for measuring image coordinates?
 - Pixel units
- Thus, we need to introduce *scaling factors* for mapping from world to pixel units


$$\begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} = \begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

calibration
matrix K

Canonical
projection
matrix $[I \mid \mathbf{0}]$

$$P = K[I \mid \mathbf{0}]$$

Intrinsic parameters: Scaling factors




Camera sensor

m_x pixels/m in horizontal direction,
 m_y pixels/m in vertical direction

Pixel size (m): $\frac{1}{m_x} \times \frac{1}{m_y}$

Scaling factors	Calibration matrix K in metric units	Calibration matrix K in pixel units
$\begin{bmatrix} m_x & 0 & 0 \\ 0 & m_y & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$\begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix}$	$= \begin{bmatrix} \alpha_x & 0 & \beta_x \\ 0 & \alpha_y & \beta_y \\ 0 & 0 & 1 \end{bmatrix}$
pixels/m	m	pixels

Extrinsic Parameters



In general, the camera coordinate frame will be related to the world coordinate frame by a rotation and a translation.

In non-homogeneous coordinates, the transformation from **world** to normalized **camera** coordinate system is given by:

$$\tilde{X}_{cam} = R(\tilde{X} - \tilde{C}) = R\tilde{X} + t$$


coords. of point in normalized camera frame

3x3 rotation matrix

coords. of a point in world frame

coords. of camera center in world frame

Extrinsic Parameters



In *non-homogeneous* coordinates:

$$\tilde{X}_{\text{cam}} = R\tilde{X} + t$$

In *homogeneous* coordinates:


$$X_{\text{cam}} = \begin{bmatrix} R & t \\ \mathbf{0}^T & 1 \end{bmatrix} X$$

3D transformation
matrix (4 x 4)

Transformation from normalized 3D coordinates to pixel image coordinates:

$$x \cong K[I|\mathbf{0}]X_{\text{cam}}$$

Extrinsic Parameters



$$x \cong K[I|0] \begin{bmatrix} R & t \\ 0^T & 1 \end{bmatrix} X$$

Simplifying:

$$x \cong K[R|t]X \quad t = -R\tilde{C}$$



Camera calibration

$$\boldsymbol{x} \cong \boldsymbol{K}[\boldsymbol{R} \quad \boldsymbol{t}]\boldsymbol{X}$$

$$\begin{pmatrix} x \\ y \\ 1 \end{pmatrix} \cong \begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

Camera calibration

Given n points with known 3D coordinates X_i and known image projections x_i , estimate the camera parameters



Camera calibration

Given n points with known 3D coordinates X_i and known image projections x_i , estimate the camera parameters



Camera calibration: Linear method

$$\mathbf{x}_i \cong \mathbf{P} \mathbf{X}_i \quad \mathbf{P}_j = \begin{pmatrix} p_{j1} \\ p_{j2} \\ p_{j3} \\ p_{j4} \end{pmatrix} \quad \begin{pmatrix} x_i \\ y_i \\ 1 \end{pmatrix} \cong \begin{pmatrix} \mathbf{X}_i^T \mathbf{P}_1 \\ \mathbf{X}_i^T \mathbf{P}_2 \\ \mathbf{X}_i^T \mathbf{P}_3 \end{pmatrix} \quad \begin{aligned} \mathbf{X}_i^T \mathbf{P}_1 - x_i \mathbf{X}_i^T \mathbf{P}_3 &= 0 \\ \mathbf{X}_i^T \mathbf{P}_2 - y_i \mathbf{X}_i^T \mathbf{P}_3 &= 0 \end{aligned}$$

$$\begin{bmatrix} \mathbf{X}_i^T & 0 & -x_i \mathbf{X}_i^T \\ 0 & \mathbf{X}_i^T & -y_i \mathbf{X}_i^T \end{bmatrix} \begin{pmatrix} \mathbf{P}_1 \\ \mathbf{P}_2 \\ \mathbf{P}_3 \end{pmatrix} = 0$$

One match gives **two** linearly independent constraints

Camera calibration: Linear method

Final linear system:

$$\begin{bmatrix} \mathbf{0}^T & \mathbf{X}_1^T & -y_1 \mathbf{X}_1^T \\ \mathbf{X}_1^T & \mathbf{0}^T & -x_1 \mathbf{X}_1^T \\ \dots & \dots & \dots \\ \mathbf{0}^T & \mathbf{X}_n^T & -y_n \mathbf{X}_n^T \\ \mathbf{X}_n^T & \mathbf{0}^T & -x_n \mathbf{X}_n^T \end{bmatrix} \begin{pmatrix} \mathbf{P}_1 \\ \mathbf{P}_2 \\ \mathbf{P}_3 \end{pmatrix} = 0 \quad \mathbf{A}\mathbf{p} = 0$$

- One 2D/3D correspondence gives two linearly independent equations
 - The projection matrix has 11 degrees of freedom
 - 6 correspondences needed for a minimal solution
- Homogeneous least squares: find \mathbf{p} minimizing $\|\mathbf{A}\mathbf{p}\|^2$
 - Solution is eigenvector of $\mathbf{A}^T\mathbf{A}$ corresponding to smallest eigenvalue

Camera calibration: Linear vs. nonlinear

Linear calibration is easy to formulate and solve, but it doesn't directly tell us the camera parameters

$$\begin{pmatrix} x \\ y \\ 1 \end{pmatrix} \cong \begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

$$\text{vs.} \quad \boldsymbol{x} \cong \boldsymbol{K}[\boldsymbol{R} \mid \boldsymbol{t}]\boldsymbol{X}$$

Camera calibration: Linear vs. nonlinear

In practice, *non-linear* methods are preferred

- Write down objective function in terms of intrinsic and extrinsic parameters, as sum of squared distances between measured 2D points and estimated projections of 3D points:

$$\sum_i \|\text{proj}(\mathbf{K}[\mathbf{R} \mid \mathbf{t}]\mathbf{X}_i) - \mathbf{x}_i\|_2^2$$

- Can include radial distortion and constraints such as known focal length, orthogonality, visibility of points
- Minimize error using non-linear optimization package
- Can initialize solution with output of linear method (perform QR decomposition to get \mathbf{K} and \mathbf{R} from \mathbf{P})